

# Task\_5\_prodigy\_infotech\_internship

June 4, 2024

PRODIGY INFOTECH DATA SCIENCE INTERN

#TASK 5

TASK OVERVIEW: Analyze traffic accident data to identify patterns related to road conditions, weather, and time of day. Visualize accident hotspots and contributing factors.

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[ ]: df = pd.read_csv("/content/traffic accident in india.csv")
```

## DATA PREPROCESSING

```
[ ]: df.head()
```

```
[ ]:      Time Day_of_week Age_band_of_driver Sex_of_driver Educational_level \
0  17:02:00      Monday          18-30          Male  Above high school
1  17:02:00      Monday          31-50          Male  Junior high school
2  17:02:00      Monday          18-30          Male  Junior high school
3  01:06:00      Sunday          18-30          Male  Junior high school
4  01:06:00      Sunday          18-30          Male  Junior high school
```

```
      Vehicle_driver_relation Driving_experience      Type_of_vehicle \
0              Employee          1-2yr          Automobile
1              Employee      Above 10yr  Public (> 45 seats)
2              Employee          1-2yr  Lorry (41?100Q)
3              Employee          5-10yr  Public (> 45 seats)
4              Employee          2-5yr              NaN
```

```
      Owner_of_vehicle Service_year_of_vehicle ... Vehicle_movement \
0              Owner      Above 10yr ...  Going straight
1              Owner          5-10yrs ...  Going straight
2              Owner              NaN ...  Going straight
3      Governmental              NaN ...  Going straight
4              Owner          5-10yrs ...  Going straight
```

	Casualty_class	Sex_of_casualty	Age_band_of_casualty	Casualty_severity	\
0	na	na	na	na	
1	na	na	na	na	
2	Driver or rider	Male	31-50	3	
3	Pedestrian	Female	18-30	3	
4	na	na	na	na	

	Work_of_casualty	Fitness_of_casualty	Pedestrian_movement	\
0	NaN	NaN	Not a Pedestrian	
1	NaN	NaN	Not a Pedestrian	
2	Driver	NaN	Not a Pedestrian	
3	Driver	Normal	Not a Pedestrian	
4	NaN	NaN	Not a Pedestrian	

	Cause_of_accident	Accident_severity
0	Moving Backward	Slight Injury
1	Overtaking	Slight Injury
2	Changing lane to the left	Serious Injury
3	Changing lane to the right	Slight Injury
4	Overtaking	Slight Injury

[5 rows x 32 columns]

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12316 entries, 0 to 12315
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Time                                  12316 non-null  object
1   Day_of_week                           12316 non-null  object
2   Age_band_of_driver                    12316 non-null  object
3   Sex_of_driver                         12316 non-null  object
4   Educational_level                     11575 non-null  object
5   Vehicle_driver_relation                11737 non-null  object
6   Driving_experience                     11487 non-null  object
7   Type_of_vehicle                       11366 non-null  object
8   Owner_of_vehicle                      11834 non-null  object
9   Service_year_of_vehicle               8388 non-null   object
10  Defect_of_vehicle                     7889 non-null   object
11  Area_accident_occured                 12077 non-null  object
12  Lanes_or_Medians                      11931 non-null  object
13  Road_allignment                       12174 non-null  object
14  Types_of_Junction                     11429 non-null  object
15  Road_surface_type                     12144 non-null  object
16  Road_surface_conditions                12316 non-null  object
```

```

17 Light_conditions          12316 non-null object
18 Weather_conditions        12316 non-null object
19 Type_of_collision          12161 non-null object
20 Number_of_vehicles_involved 12316 non-null int64
21 Number_of_casualties       12316 non-null int64
22 Vehicle_movement          12008 non-null object
23 Casualty_class             12316 non-null object
24 Sex_of_casualty            12316 non-null object
25 Age_band_of_casualty       12316 non-null object
26 Casualty_severity          12316 non-null object
27 Work_of_casualty           9118 non-null object
28 Fitness_of_casualty        9681 non-null object
29 Pedestrian_movement        12316 non-null object
30 Cause_of_accident          12316 non-null object
31 Accident_severity          12316 non-null object
dtypes: int64(2), object(30)
memory usage: 3.0+ MB

```

```
[ ]: df.describe()
```

```

[ ]:      Number_of_vehicles_involved  Number_of_casualties
count                12316.000000                12316.000000
mean                   2.040679                   1.548149
std                    0.688790                   1.007179
min                    1.000000                   1.000000
25%                    2.000000                   1.000000
50%                    2.000000                   1.000000
75%                    2.000000                   2.000000
max                    7.000000                   8.000000

```

```
[ ]: df.columns
```

```

[ ]: Index(['Time', 'Day_of_week', 'Age_band_of_driver', 'Sex_of_driver',
          'Educational_level', 'Vehicle_driver_relation', 'Driving_experience',
          'Type_of_vehicle', 'Owner_of_vehicle', 'Service_year_of_vehicle',
          'Defect_of_vehicle', 'Area_accident_occured', 'Lanes_or_Medians',
          'Road_allignment', 'Types_of_Junction', 'Road_surface_type',
          'Road_surface_conditions', 'Light_conditions', 'Weather_conditions',
          'Type_of_collision', 'Number_of_vehicles_involved',
          'Number_of_casualties', 'Vehicle_movement', 'Casualty_class',
          'Sex_of_casualty', 'Age_band_of_casualty', 'Casualty_severity',
          'Work_of_casualty', 'Fitness_of_casualty', 'Pedestrian_movement',
          'Cause_of_accident', 'Accident_severity'],
          dtype='object')

```

```
[ ]: df.nunique()
```

```
[ ]: Time 1074
    Day_of_week 7
    Age_band_of_driver 5
    Sex_of_driver 3
    Educational_level 7
    Vehicle_driver_relation 4
    Driving_experience 7
    Type_of_vehicle 17
    Owner_of_vehicle 4
    Service_year_of_vehicle 6
    Defect_of_vehicle 3
    Area_accident_occured 14
    Lanes_or_Medians 7
    Road_allignment 9
    Types_of_Junction 8
    Road_surface_type 5
    Road_surface_conditions 4
    Light_conditions 4
    Weather_conditions 9
    Type_of_collision 10
    Number_of_vehicles_involved 6
    Number_of_casualties 8
    Vehicle_movement 13
    Casualty_class 4
    Sex_of_casualty 3
    Age_band_of_casualty 6
    Casualty_severity 4
    Work_of_casualty 7
    Fitness_of_casualty 5
    Pedestrian_movement 9
    Cause_of_accident 20
    Accident_severity 3
    dtype: int64
```

```
[ ]: df.isnull().sum()
```

```
[ ]: Time 0
    Day_of_week 0
    Age_band_of_driver 0
    Sex_of_driver 0
    Educational_level 741
    Vehicle_driver_relation 579
    Driving_experience 829
    Type_of_vehicle 950
    Owner_of_vehicle 482
    Service_year_of_vehicle 3928
    Defect_of_vehicle 4427
```

Area_accident_occured	239
Lanes_or_Medians	385
Road_allignment	142
Types_of_Junction	887
Road_surface_type	172
Road_surface_conditions	0
Light_conditions	0
Weather_conditions	0
Type_of_collision	155
Number_of_vehicles_involved	0
Number_of_casualties	0
Vehicle_movement	308
Casualty_class	0
Sex_of_casualty	0
Age_band_of_casualty	0
Casualty_severity	0
Work_of_casualty	3198
Fitness_of_casualty	2635
Pedestrian_movement	0
Cause_of_accident	0
Accident_severity	0

dtype: int64

```
[ ]: # fill missing values with mean column values
df['Driving_experience'].fillna(df['Driving_experience'].mode()[0],
    inplace=True)
df['Age_band_of_driver'].fillna(df['Age_band_of_driver'].mode()[0],
    inplace=True)
df['Type_of_vehicle'].fillna(df['Type_of_vehicle'].mode()[0], inplace=True)
df['Area_accident_occured'].fillna(df['Area_accident_occured'].mode()[0],
    inplace=True)
df['Road_allignment'].fillna(df['Road_allignment'].mode()[0], inplace=True)
df['Type_of_collision'].fillna(df['Type_of_collision'].mode()[0], inplace=True)
df['Vehicle_movement'].fillna(df['Vehicle_movement'].mode()[0], inplace=True)
df['Lanes_or_Medians'].fillna(df['Lanes_or_Medians'].mode()[0], inplace=True)
df['Types_of_Junction'].fillna(df['Types_of_Junction'].mode()[0], inplace=True)
```

```
[ ]: df.isnull().sum()
```

Age_band_of_driver	0
Driving_experience	0
Type_of_vehicle	0
Area_accident_occured	0
Lanes_or_Medians	0
Road_allignment	0
Types_of_Junction	0
Road_surface_conditions	0

```

Light_conditions      0
Weather_conditions    0
Type_of_collision     0
Number_of_vehicles_involved  0
Number_of_casualties  0
Vehicle_movement      0
Casualty_class        0
Age_band_of_casualty  0
Pedestrian_movement  0
Cause_of_accident     0
Accident_severity     0
TimeInSeconds         0
dtype: int64

```

```
[ ]: df.duplicated()
```

```

[ ]: 0      False
      1      False
      2      False
      3      False
      4      False
      ...
     12311   False
     12312   False
     12313   False
     12314   False
     12315   False
      Length: 12316, dtype: bool

```

```
[ ]: df.dtypes.value_counts()
```

```

[ ]: object    30
      int64     2
      Name: count, dtype: int64

```

```
[ ]: # convert the 'Date' column to datetime format
      df['Time'] = pd.to_datetime(df['Time'])
```

<ipython-input-16-e69021690a5a>:2: UserWarning: Could not infer format, so each element will be parsed individually, falling back to `dateutil`. To ensure parsing is consistent and as-expected, please specify a format.

```
df['Time'] = pd.to_datetime(df['Time'])
```

## EDA

```
[ ]: plt.figure(figsize=(10, 5))
      plt.subplot(1, 2, 1)
```

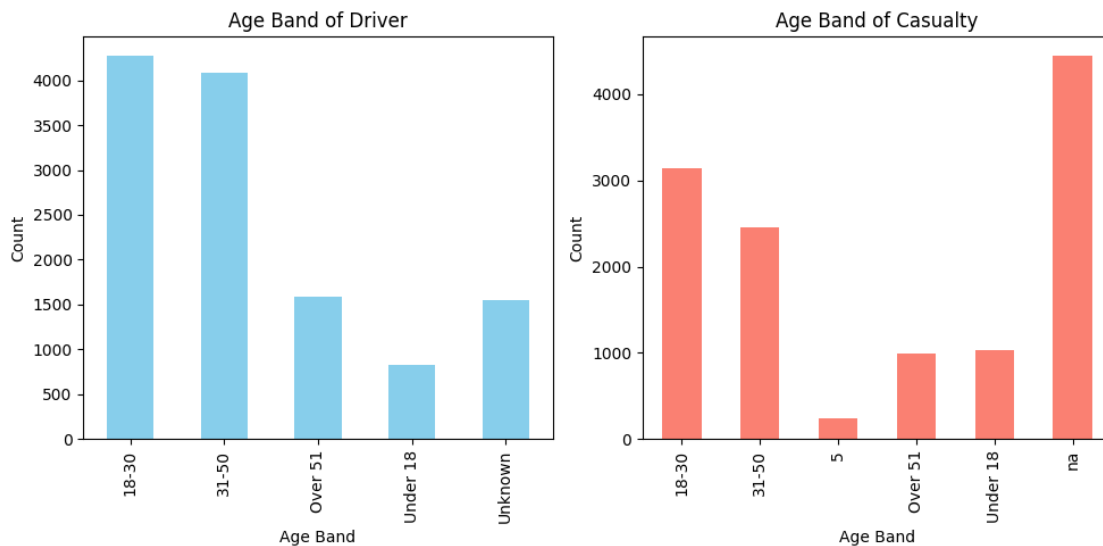
```

df['Age_band_of_driver'].value_counts().sort_index().plot(kind='bar',
    ↪color='skyblue')
plt.title('Age Band of Driver')
plt.xlabel('Age Band')
plt.ylabel('Count')

# 'Age_band_of_casualty'
plt.subplot(1, 2, 2)
df['Age_band_of_casualty'].value_counts().sort_index().plot(kind='bar',
    ↪color='salmon')
plt.title('Age Band of Casualty')
plt.xlabel('Age Band')
plt.ylabel('Count')

plt.tight_layout()
plt.show()


```



```

[ ]: df.hist(layout=(1,6), figsize=(30,5))
plt.show()

```




```
[ ]: df['Number_of_casualties'].value_counts()
```


```
[ ]: Number_of_casualties
1      8397
2      2290
3       909
4       394
5       207
6        89
7        22
8         8
Name: count, dtype: int64
```

```
[ ]: plt.figure(figsize=(10,7))
sns.boxplot(data=df, y='Number_of_vehicles_involved',
            x='Number_of_casualties',color="lightgreen")
plt.show()
```






```
[ ]: sns.boxplot(data=df, y='Number_of_casualties')  
plt.show()
```




```
[ ]: sns.boxplot(data=df, y='Number_of_vehicles_involved')  
plt.show()
```




```
[ ]: df['Number_of_vehicles_involved']
```

```
[ ]: 0      2
      1      2
      2      2
      3      2
      4      2
      ..
12311    2
12312    2
12313    1
12314    2
12315    2
Name: Number_of_vehicles_involved, Length: 12316, dtype: int64
```

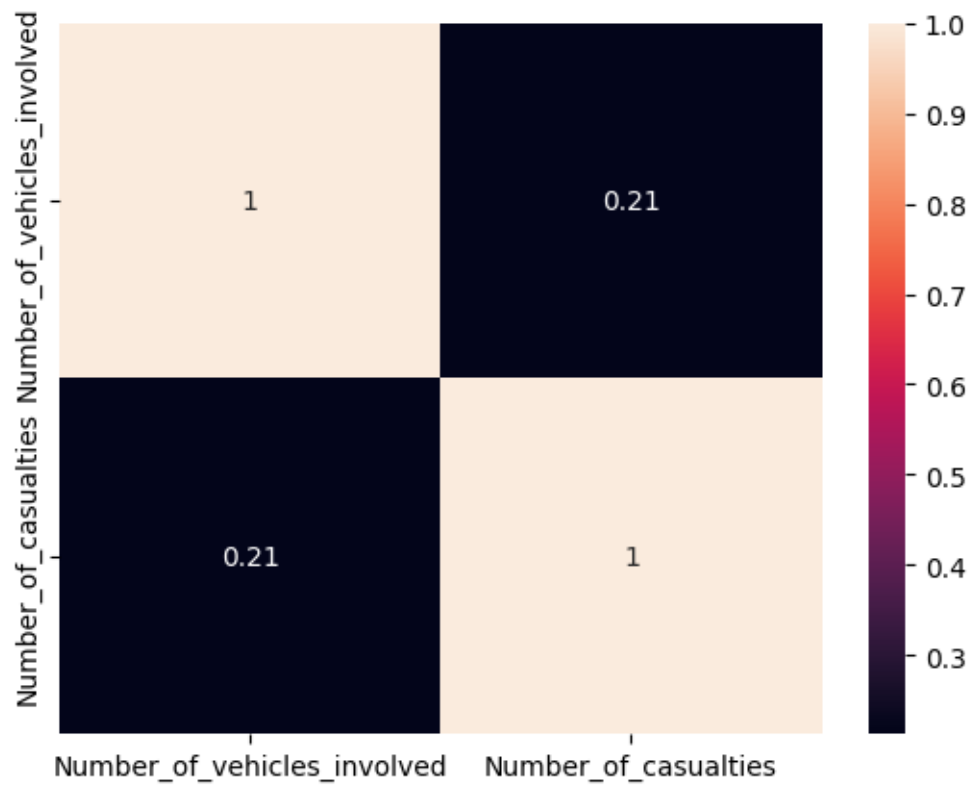
```
[ ]: sns.scatterplot(x=df['Number_of_vehicles_involved'],
                    ↪y=df['Number_of_casualties'])
plt.show()
```




```
[ ]: sns.pairplot(df[['Number_of_vehicles_involved', 'Number_of_casualties']])  
plt.show()
```



```
[ ]: correlation_matrix = df[['Number_of_vehicles_involved', 'Number_of_casualties']].
    ↪corr()
sns.heatmap(correlation_matrix, annot=True)
plt.show()
```



```
[ ]: plt.figure(figsize=(10,7))
plt.pie(x=df['Accident_severity'].value_counts().values,
        labels=df['Accident_severity'].value_counts().index,
        autopct='%2.2f%%')
plt.show()
```



```
[ ]: # creating a facet grid with columns as survived=0 and survived=1
grid = sns.FacetGrid(data=df, col='Accident_severity', height=4, aspect=1,
    ↳sharey=False)
# mapping bar plot and the data on to the grid
grid.map(sns.countplot, 'Number_of_vehicles_involved', palette=['black',
    ↳'brown', 'orange'])
plt.show()
```

/usr/local/lib/python3.10/dist-packages/seaborn/axisgrid.py:718: UserWarning:  
Using the countplot function without specifying `order` is likely to produce an  
incorrect plot.

warnings.warn(warning)

/usr/local/lib/python3.10/dist-packages/seaborn/axisgrid.py:854: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in  
v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same

effect.

```
func(*plot_args, **plot_kwargs)
/usr/local/lib/python3.10/dist-packages/seaborn/axisgrid.py:854: UserWarning:
The palette list has fewer values (3) than needed (6) and will cycle, which may
produce an uninterpretable plot.
```

```
func(*plot_args, **plot_kwargs)
/usr/local/lib/python3.10/dist-packages/seaborn/axisgrid.py:854: FutureWarning:
```


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
func(*plot_args, **plot_kwargs)
/usr/local/lib/python3.10/dist-packages/seaborn/axisgrid.py:854: UserWarning:
The palette list has fewer values (3) than needed (5) and will cycle, which may
produce an uninterpretable plot.
```

```
func(*plot_args, **plot_kwargs)
/usr/local/lib/python3.10/dist-packages/seaborn/axisgrid.py:854: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
func(*plot_args, **plot_kwargs)
```




```
[ ]: df['TimeInSeconds'] = df['Time'].apply(lambda x: x.hour * 3600 + x.minute * 60 + x.second)
```

```
plt.figure(figsize=(10,6))
plt.hist(df['TimeInSeconds'], bins=24, color='blue', edgecolor='black')
plt.title('Traffic Accident Time Distribution')
plt.xlabel('Time (Hours)')
plt.ylabel('Number of Accidents')
```




```
plt.xticks(range(0, 24 * 3600, 3600), range(24))  
plt.show()
```



```
[ ]: plt.figure(figsize=[25,15])  
     sns.heatmap(df.corr(),annot=True)
```

```
[ ]: <Axes: >
```



## FEATURE ENGINEERING

```
[ ]: # dropping columns that can cause imbalance while imputation
lists=['Vehicle_driver_relation', 'Work_of_casualty',
      ↪ 'Fitness_of_casualty', 'Day_of_week', 'Casualty_severity', 'Time', 'Sex_of_driver', 'Educational_level',
      ↪ 'Road_surface_type', 'Sex_of_casualty']
df.drop(columns = lists, inplace=True)
```

```
[ ]: df.columns
```

```
[ ]: Index(['Age_band_of_driver', 'Driving_experience', 'Type_of_vehicle',
          'Area_accident_occured', 'Lanes_or_Medians', 'Road_alignment',
          'Types_of_Junction', 'Road_surface_conditions', 'Light_conditions',
          'Weather_conditions', 'Type_of_collision',
          'Number_of_vehicles_involved', 'Number_of_casualties',
          'Vehicle_movement', 'Casualty_class', 'Age_band_of_casualty',
          'Pedestrian_movement', 'Cause_of_accident', 'Accident_severity',
          'TimeInSeconds'],
          dtype='object')
```

```
[ ]: from sklearn.preprocessing import LabelEncoder #or one hot encoder
LE = LabelEncoder()
```

```
df=df.apply(LE.fit_transform)
```

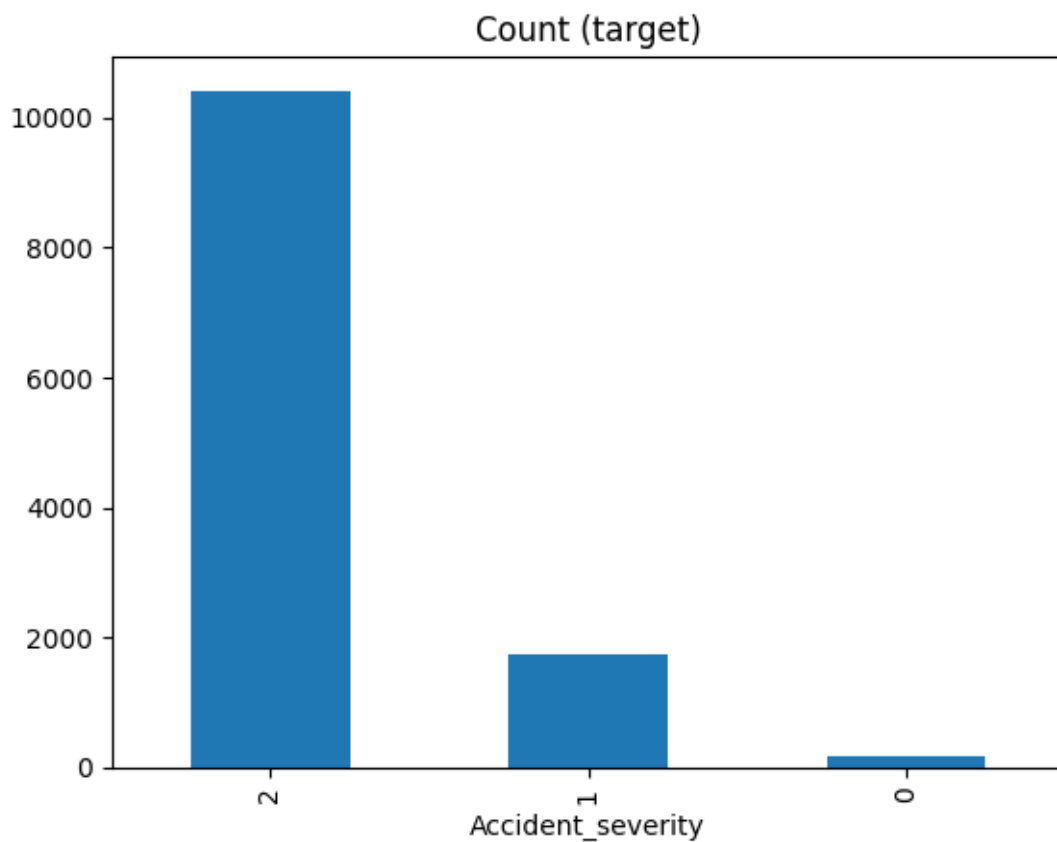
```
[ ]: target_count = df['Accident_severity'].value_counts()
print('Class 0:', target_count[0])
print('Class 1:', target_count[1])
print('Proportion:', round(target_count[0] / target_count[1], 2), ': 1')

target_count.plot(kind='bar', title='Count (target)');
```

Class 0: 158

Class 1: 1743

Proportion: 0.09 : 1



## UPSAMPLING

```
[ ]: from sklearn.model_selection import train_test_split
```

```
[ ]: x = df.drop('Accident_severity', axis=1)
y = df['Accident_severity']
```

```
xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size=0.3,
        random_state=42)
print(xtrain.shape, xtest.shape, ytrain.shape, ytest.shape)
```

```
(8621, 19) (3695, 19) (8621,) (3695,)
```

```
[ ]: from collections import Counter
      from imblearn.over_sampling import SMOTE
```

```
[ ]: # upsampling using smote

counter = Counter(ytrain)

print("=====")

for k,v in counter.items():
    per = 100*v/len(ytrain)
    print(f"Class= {k}, n={v} ({per:.2f}%)")

oversample = SMOTE()
xtrain, ytrain = oversample.fit_resample(xtrain, ytrain)

counter = Counter(ytrain)

print("=====")

for k,v in counter.items():
    per = 100*v/len(ytrain)
    print(f"Class= {k}, n={v} ({per:.2f}%)")

print("=====")

print("Upsampled data shape: ", xtrain.shape, ytrain.shape)
```

```
=====
Class= 2, n=7324 (84.96%)
Class= 1, n=1191 (13.82%)
Class= 0, n=106 (1.23%)
=====
Class= 2, n=7324 (33.33%)
Class= 1, n=7324 (33.33%)
Class= 0, n=7324 (33.33%)
=====
Upsampled data shape: (21972, 19) (21972,)
```

## SPLITTING AND MODEL TRAINING

```
[ ]: x=df.drop(columns=["Accident_severity"])
      y=df["Accident_severity"]
```

```
[ ]: from sklearn.linear_model import LogisticRegression
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.svm import SVC
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.naive_bayes import GaussianNB
      from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier,
      ↪ GradientBoostingClassifier
```

```
[ ]: # models={"LogisticRegression":LogisticRegression(),
      #        "DecisionTreeClassifier":DecisionTreeClassifier(),
      #        "SVM":SVC(),
      #        "KNeighborsClassifier":KNeighborsClassifier(),
      #        "GNB":GaussianNB(),
      #        "RandomForestClassifier":RandomForestClassifier(),
      #        "AdaBoostClassifier":AdaBoostClassifier(),
      #        "GradientBoostingClassifier":GradientBoostingClassifier(),
      #        }
      from sklearn.multiclass import OneVsRestClassifier
      models = {
          "LogisticRegression": OneVsRestClassifier(LogisticRegression(C=1.0,
          ↪ solver='liblinear', penalty='l2')),
          "DecisionTreeClassifier":
          ↪ OneVsRestClassifier(DecisionTreeClassifier(criterion='entropy', max_depth=5,
          ↪ min_samples_split=5)),
          "KNeighborsClassifier":
          ↪ OneVsRestClassifier(KNeighborsClassifier(n_neighbors=5, weights='distance',
          ↪ p=2)),
          "RandomForestClassifier":
          ↪ OneVsRestClassifier(RandomForestClassifier(n_estimators=100,
          ↪ criterion='entropy', max_depth=8)),
          "AdaBoostClassifier":
          ↪ OneVsRestClassifier(AdaBoostClassifier(n_estimators=200, learning_rate=0.1)),
          "GradientBoostingClassifier":
          ↪ OneVsRestClassifier(GradientBoostingClassifier(n_estimators=100,
          ↪ learning_rate=0.1, max_depth=5))
      }
```

## MODEL EVALUATION

```
[ ]: from sklearn.pipeline import Pipeline
      from sklearn.metrics import classification_report, accuracy_score
      from sklearn.preprocessing import MinMaxScaler, StandardScaler
      from sklearn.metrics import roc_auc_score
      from sklearn.metrics import roc_curve
```

```
[ ]: # models,x,y,scaleFlag=0,1,2
def modelAccuracy(models,x,y,scaleFlag):
    #train/Test
    xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.3,random_state=0)
    acc_result={}
    for name,model in models.items():
        #pipeline
        #1.Transformer -> 2.Model

    ↪print("=====",name,"=====")
        if(scaleFlag==1):
            ↪model_pipeline=Pipeline([('MinMaxScler',MinMaxScaler()),('model',model)])
            elif(scaleFlag==2):
                ↪model_pipeline=Pipeline([('StandardScaler',StandardScaler()),('model',model)])
            else:
                model_pipeline=Pipeline([('model',model)])
                #training/testing on model pipeline
                model_fit=model_pipeline.fit(xtrain,ytrain)
                ypred=model_fit.predict(xtest)
                ypred_proba=model_fit.predict_proba(xtest)
                ypred_proba = np.nan_to_num(ypred_proba, nan=1/3)

                print(classification_report(ytest,ypred))
                auc = roc_auc_score(ytest, ypred_proba, multi_class='ovr',
    ↪average='macro')
                print("macro-AUC:",auc)
                auc = roc_auc_score(ytest, ypred_proba, multi_class='ovr',
    ↪average='weighted')
                print("weighted-AUC:",auc,"\n")
                plt_auc(ytest,ypred_proba)

acc=modelAccuracy(models,x,y,1)
```

===== LogisticRegression =====

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/\_classification.py:1344:  
 UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
 0.0 in labels with no predicted samples. Use `zero\_division` parameter to  
 control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))  
 /usr/local/lib/python3.10/dist-packages/sklearn/metrics/\_classification.py:1344:  
 UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
 0.0 in labels with no predicted samples. Use `zero\_division` parameter to  
 control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))


```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	51
1	0.00	0.00	0.00	537
2	0.84	1.00	0.91	3107
accuracy			0.84	3695
macro avg	0.28	0.33	0.30	3695
weighted avg	0.71	0.84	0.77	3695

```
macro-AUC: 0.6213067642957238
```

```
weighted-AUC: 0.5778686446039931
```



```
===== DecisionTreeClassifier
```

```

=====
              precision    recall  f1-score   support


     0       0.83         0.10         0.18         51
     1       0.53         0.02         0.04        537
     2       0.84         1.00         0.91       3107

 accuracy          0.84         3695
 macro avg       0.73         0.37         0.38         3695
 weighted avg    0.80         0.84         0.78         3695

```

macro-AUC: 0.6792863278673856

weighted-AUC: 0.620070910304444



```

===== KNeighborsClassifier
=====
              precision    recall  f1-score   support

     0       0.00         0.00         0.00         51


```



1	0.15	0.02	0.04	537
2	0.84	0.98	0.90	3107
accuracy			0.83	3695
macro avg	0.33	0.33	0.31	3695
weighted avg	0.73	0.83	0.77	3695

macro-AUC: 0.5631938253966587

weighted-AUC: 0.558606823502534



```
===== RandomForestClassifier
=====
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
```

0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```


/usr/local/lib/python3.10/dist-packages/sklearn/metrics/\_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero\_division` parameter to  
control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	51
1	1.00	0.00	0.00	537
2	0.84	1.00	0.91	3107
accuracy			0.84	3695
macro avg	0.61	0.33	0.31	3695
weighted avg	0.85	0.84	0.77	3695

macro-AUC: 0.7285297276891907

weighted-AUC: 0.6661154702181663



```

===== AdaBoostClassifier =====

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))


      precision    recall  f1-score   support

0         0.00      0.00      0.00         51
1         0.00      0.00      0.00        537
2         0.84      1.00      0.91       3107

 accuracy          0.84       3695
 macro avg         0.28      0.33      0.30       3695
weighted avg         0.71      0.84      0.77       3695

macro-AUC: 0.6680061303289714
weighted-AUC: 0.6169814922826339

```



===== GradientBoostingClassifier

=====

	precision	recall	f1-score	support
0	0.70	0.14	0.23	51
1	0.72	0.09	0.16	537
2	0.85	0.99	0.92	3107
accuracy			0.85	3695
macro avg	0.76	0.41	0.44	3695
weighted avg	0.83	0.85	0.80	3695

macro-AUC: 0.7843804166417275

weighted-AUC: 0.7213491590919057

